

Homework 5

Jiahong Long

18 February 2021

1. On the Frobenius norm

(a) Norm equality

We perform a column-vector decomposition of A :

$$A = \begin{bmatrix} a_{:,1} & a_{:,2} & a_{:,3} & \dots & a_{:,n} \end{bmatrix}$$

By the definition of matrix multiplication, it follows immediately from this decomposition that

$$UA = \begin{bmatrix} Ua_{:,1} & Ua_{:,2} & Ua_{:,3} & \dots & Ua_{:,n} \end{bmatrix}$$

Recall that the Frobenius norm of a matrix M is defined as

$$\|M\|_F = \sqrt{\sum_{i=1}^n \sum_{j=1}^n |m_{i,j}|^2}$$

which is in turn equivalent to the sum of the Euclidean norm of its columns

$$= \sqrt{\sum_{i=1}^n \|m_{:,i}\|_2^2}$$

Given that U is $n \times n$ orthogonal, we know that the norm of a vector in \mathbb{R}^n is the same as its image under U . Then it immediately follows that

$$\|m_{:,i}\|_2 = \|Um_{:,i}\|_2 \quad \rightarrow \quad \|m_{:,i}\|_2^2 = \|Um_{:,i}\|_2^2$$

for all $i = 1, 2, \dots, n$. Then, substituting the above into the formula for the definition of the Frobenius norm immediately yields equality:

$$\begin{aligned} \|M\|_F &= \sqrt{\sum_{i=1}^n \|m_{:,i}\|_2^2} \\ &= \sqrt{\sum_{i=1}^n \|Um_{:,i}\|_2^2} \end{aligned}$$

Then, by definition,

$$= \|UM\|_F$$

Let $M = A$, and

$$\|A\|_F = \|UA\|_F$$

as required. ■

(b) *Norm equality redux*

We begin by transposing both sides of the equation

$$\begin{aligned} A &= AV \\ A^T &= (AV)^T \\ &= V^T A^T \end{aligned}$$

Define $A' = A^T$, and

$$A' = V^T A'$$

Since V is orthogonal, we know that V^T is also orthogonal. The proof at this point proceeds in an identical fashion to part (a); substitute A' for A and V^T for U . ■

2. *Frobenius norm and singular values*

(a) *Verification*

We begin with the singular value decomposition of A :

$$\begin{aligned} A &= U\Sigma V^T \\ \|A\|_F &= \|U\Sigma V^T\|_F \end{aligned}$$

Recall that U and V^T are both orthogonal matrices in \mathbb{R}^n . It follows, therefore, that, applying the results of (1a-b), we have

$$\begin{aligned} \|U\Sigma V^T\|_F &= \|\Sigma V^T\|_F \\ &= \|\Sigma\|_F \end{aligned}$$

Using the definition of the Frobenius norm, we rewrite this as a function of entries in Σ as

$$\|\Sigma\|_F = \sqrt{\sum_{i=1}^n \sum_{j=1}^n |\Sigma_{i,j}|^2}$$

Since Σ is nonzero only on the main diagonal, where the entries are the singular values of A , this is given by

$$= \sqrt{\sum_{i=1}^n |\sigma_i|^2}$$

Singular values are all positive, so we may drop the absolute value to obtain

$$\|A\|_F = \sqrt{\sum_{i=1}^n \sigma_i^2}$$

which is as required. ■

(b) *Condition number*

We begin by finding $\|A^{-1}\|_F$ for arbitrary $A \in \mathbb{R}^n$. To begin, note that the “SVD” of A^{-1} is given by the expression

$$A^{-1} = V\Sigma^{-1}U^T$$

where Σ^{-1} is the element-wise reciprocal of Σ . Since the Frobenius norm of matrices is unchanged through multiplication with orthogonal matrices, we may state

$$\|A^{-1}\|_F = \|V\Sigma^{-1}U^T\|_F = \|\Sigma^{-1}\|_F$$

and since Σ^{-1} has nonzero values strictly on the diagonal equal to the reciprocal of the (necessarily positive) singular values, this is in turn given by

$$\|A^{-1}\|_F = \sqrt{\sum_{i=1}^n \frac{1}{\sigma_i^2}}$$

Recall that the condition number is defined as $\kappa(A) = \|A\| \|A^{-1}\|$. Under the Frobenius norm, this is given by

$$\|A\|_F \|A^{-1}\|_F$$

and with the above results, we have

$$\sqrt{\left(\sum_{i=0}^n \sigma_i^2\right) \left(\sum_{i=0}^n \frac{1}{\sigma_i^2}\right)}$$

3. $n \times m$ matrix(a) *SVD calculation*

We know that

$$A = U\Sigma V^T$$

It follows immediately that

$$A^T = V\Sigma U^T$$

and as such

$$\begin{aligned} A^T A &= V\Sigma U^T (U\Sigma V^T) \\ &= V\Sigma^T \cancel{U^T U}^{\nearrow I_n} \Sigma V^T \\ &= V\Sigma^T \Sigma V^T \\ &= V \underbrace{\begin{bmatrix} \sigma_1^2 & 0 & 0 & \dots & 0 \\ 0 & \sigma_2^2 & 0 & \dots & 0 \\ 0 & 0 & \sigma_3^2 & \dots & 0 \\ \vdots & & & \ddots & \vdots \\ 0 & 0 & 0 & \dots & \sigma_m^2 \end{bmatrix}}_{\hat{\Sigma}_2} V^T \end{aligned}$$

where the middle matrix $\hat{\Sigma}_2$ is just the reduced singular value matrix with all entries squared.

$$= \boxed{V\hat{\Sigma}_2 V^T}$$

We posit that this in fact the singular value decomposition of $A^T A$ because V and its transpose are both orthogonal in \mathbb{R}^n , and $\hat{\Sigma}_2$ is a diagonal matrix with all nonnegative entries in increasing order.

(b) *Conditioning equality*

Spectral norm proof

Let $A = U\Sigma V^T$, and per part (a) let $A^T A = V\hat{\Sigma}_2 V^T$. Recall that the spectral matrix norm of a matrix is given exactly by its largest singular value σ_1 . It therefore follows that

$$\|A\|_2 = \sigma_1$$

and since the SVD of $A^T A$ is given by $V\hat{\Sigma}_2 V^T$, the largest singular value of $A^T A$ must be the leftmost entry on the top row of $\hat{\Sigma}_2$. From the previous problem this is given as exactly σ_1^2 , and we have

$$\|A^T A\|_2 = \sigma_1^2 = (\sigma_1)^2 = \|A\|_2^2$$

as required. ■

Condition number proof

We begin by finding the condition number of A under the spectral norm. Recall that the spectral matrix norm is just the largest singular value of that matrix

$$\|A\|_2 = \sigma_1$$

It follows that the spectral norm of the matrix *inverse* must be the largest singular value of the matrix inverse, which is just the reciprocal of the smallest singular value of the matrix. Here, as A is full-rank, we therefore have

$$\|A^{-1}\|_2 = \frac{1}{\sigma_n}$$

It follows that

$$\kappa_2(A) = \frac{\sigma_1}{\sigma_n}$$

We now find an expression for $\|(A^T A)^{-1}\|_F$:

$$\begin{aligned} (A^T A)^{-1} &= (V\hat{\Sigma}_2 V^T)^{-1} \\ &= V\hat{\Sigma}_2^{-1} V^T \end{aligned}$$

Recall that the spectral norm simply returns the largest singular value of the matrix. Then we must have

$$\|(A^T A)^{-1}\|_2 = \|V\hat{\Sigma}_2^{-1} V^T\|_2$$

Note that σ_n is nonzero since A is full-rank. Furthermore, since $\hat{\Sigma}_2$ is diagonal, its inverse is obtained by taking the reciprocal of each entry along the diagonal; it follows that the largest entry in the inverse must be the reciprocal of the smallest singular value. We thus have

$$\begin{aligned} &= \frac{1}{|\sigma_n^2|} \\ &= \frac{1}{\sigma_n^2} \end{aligned}$$

It follows immediately that

$$\kappa_2(A^T A) = \|A^T A\|_2 \|(A^T A)^{-1}\|_2 = \frac{\sigma_1^2}{\sigma_n^2}$$

and since

$$\kappa_2(A) = \frac{\sigma_1}{\sigma_n}$$

we immediately have $\kappa_2(A^T A) = \kappa_2(A)^2$ as required. ■

4. Computations

(a) Rank

Using elementary row operations, we obtain

$$\begin{bmatrix} 1 & 2 \\ 2 & 4 \\ 3 & 6 \end{bmatrix} \sim \begin{bmatrix} 1 & 2 \\ 0 & 0 \\ 3 & 6 \end{bmatrix} \sim \begin{bmatrix} 1 & 2 \\ 0 & 0 \\ 1 & 2 \end{bmatrix} \sim \begin{bmatrix} 1 & 2 \\ 0 & 0 \\ 0 & 0 \end{bmatrix}$$

Thus the rank of A is $\boxed{1}$.

(b) Size

Since the rank of A is 1, there must only be one nonzero singular value. It follows immediately that

$$\boxed{\Sigma_r : 1 \times 1 \quad \text{and} \quad U_r : 3 \times 1 \quad \text{and} \quad V_r^T : 1 \times 2 \rightarrow V_r : 2 \times 1}$$

(c) Submatrices

We know that there is just one singular value. Additionally, it's trivially true that

$$\underbrace{\begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}}_u \underbrace{\begin{bmatrix} 1 \\ 2 \end{bmatrix}^T}_v = \underbrace{\begin{bmatrix} 1 & 2 \\ 2 & 4 \\ 3 & 6 \end{bmatrix}}_A$$

It must follow that to obtain the reduced SVD of A , we need only normalise u to obtain U_r , and normalise v to obtain V_r . (We note here that in this case, Σ_r is vacuously equal to the scalar matrix containing the product of $\|u\|_2$ and $\|v\|_2$.) It follows that we immediately have

$$\boxed{U_r = \frac{\sqrt{14}}{14} \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix} \quad \text{and} \quad V_r = \frac{\sqrt{5}}{5} \begin{bmatrix} 1 \\ 2 \end{bmatrix}}$$

Furthermore, simple algebra yields that $\sigma_1 = \sqrt{70}$ and $\Sigma_r = [\sqrt{70}]$.

5. Low-rank approximation

(a) Line-by-line

The code in question is attached:

```

1  % load image
2  A = imread('street2.jpg');
3  A = rgb2gray(A);
4  B = double(A);
5
6  % compute SVD
7  size(B)
8  r = rank(B)
9  [U,S,V] = svd(B);
10
11 % approximate image
12 ranks = [1 2 4 8 16 32 64 r];
13 l = length(ranks);
14
15 for i = 1:l
16     % compute rank i approximation
17     k = ranks(i);
18     approxB = U(:,1:k)*S(1:k,1:k)*V(:,1:k)';
19     approxA = uint8(approxB);
20
21 % plot images
22 figure(1)
23 subplot(2,4,i)
24 imshow(approxA);
25 title(sprintf('rank %d approximation',k))
26 end

```

The line-by-line description is as follows (skipping over comments and empty newlines):

- i. The image is loaded to an integer array.
 - ii. The image is converted to grayscale.
 - iii. The image is cast to a floating-point representation.
 - iv. The size of the image is obtained in pixels.
 - v. The numerical rank of the image matrix is calculated.
 - vi. The full SVD of the image matrix is calculated.
 - vii. An array of proposed ranks is created.
 - viii. The number of proposed ranks is stored.
 - ix. We begin looping over each of the proposed ranks.
 - x. We get the current proposed rank k .
 - xi. The proposed rank is used to generate the rank- k approximation of B by multiplying out the truncated SVD of B .
 - xii. The low-rank approximation of B is converted to an integer array.
 - xiii. The current **Figure** window is set to 1 (so we use the same window for plotting).
 - xiv. We set the plot location to the i th subplot in the 2×4 grid of subplots.
 - xv. The image is plotted.
 - xvi. The image is given the title **rank k approximation** where k is the current rank.
- Hence we have a line-by-line description of the algorithm.

(b) *Algorithm*

The algorithm loads an image as a matrix representing the grayscale values of each pixel in the image. It then calculates seven different low-rank approximations of the matrix, and then plots the results as pixel arrays, along with the original (full-rank) grayscale array.

(c) *Original vs. approximation*

Low-rank approximations will use less storage and be more compact, since they allow us to represent “most” of the image using submatrices of U , Σ , and V . However, computing the actual full SVD in addition to the low-rank approximation is computationally intense, since (modern) digital images tend to have high pixel counts, leading to large matrices which are computationally difficult to deal with. For instance, a modestly sized 4000 pixel square image took about 20 seconds just to compute the rank on my computer, and even longer for the full SVD.

(d) *Reasonable k*

Heuristically (and informed, to some extent, by data science), each of the singular values σ_i can be correlated with how much of the information entropy (in data science, “variance”) is “captured” by adding the i th dimension. It stands to reason that a reasonable k can be found by fixing some proportion of the total “variance” (ie. sum of all the singular values), and finding the k where the sum of the first k th singular values meets or exceeds that proportion.

Alternatively, for a more quantitative (albeit computationally heavy) approach, one can fix some ε , and run logarithmic binary search on $k = 1, 2, 3, \dots, \text{rank}\{A\}$ to find the value of k for which all entries in the rank- k approximation of A are within ε of A . We note that logarithm binary search is a valid search method: as informed by principal component analysis in data science, scree plots of singular values necessarily decrease in a logarithmic fashion.